


Week 1 - Draw the Path: Portfolio Sitemap + Toolkit


Sitemap sketch



Change after the pressure test:

I will place my Multi-Agent/RAG project first in Selected Work and make its end-to-end evidence explicit, including architecture, backend integration, evaluation, and deployment.

Configured Claude Project



Project instructions

Set project instructions

Provide Claude with relevant instructions and information for chats within AI Engineering Portfolio Build. This will work alongside your [profile instructions](#) and the selected style in a chat.

My proof statement is: "I build and ship AI systems end-to-end, from model and retrieval design to backend APIs, evaluation, and production deployment."

My primary audience is recruiters, hiring managers, and technical collaborators looking for AI engineering experience.

The one action I want visitors to take is to contact me for AI engineering opportunities or collaborations.

Act as a tutor, not just a generator. Challenge weak assumptions, ask whether each section supports the proof statement and the one action, and explain why something should be added, removed, or simplified.

Keep recommendations concise, practical, and focused on evidence rather than decoration.

CancelSave instructions

Pressure-test prompt + beginning of output

AI Engineering Portfolio Build / Portfolio sitemap pressure-test against proof statement ▾

Pressure-test my portfolio sitemap against my proof statement and my one action.
My sitemap:

- Hero — clearly state my AI engineering claim
- Selected Work — TensorFlow open-source contribution, Webometrics ranking research, and Multi-Agent/RAG systems
- About — short background and AI/backend engineering focus
- Contact — clear way to contact me

My proof statement:

"I build and ship AI systems end-to-end, from model and retrieval design to backend APIs, evaluation, and production deployment."

My one action:

Contact me for AI engineering opportunities or collaborations.

Act as my tutor and pressure-test this rather than simply approving it. Tell me:

1. Does every section earn its place?
2. Is anything missing or unnecessary?
3. Should anything be reordered?
4. What is the single most important change I should make so the portfolio proves my claim more convincingly?

Show less

Thought for 16s

Here's the pressure test — not a rubber stamp.

1. Does every section earn its place?

Hero — Earns it, conditionally. "Clearly state my AI engineering claim" is right, but a claim alone doesn't do work. The hero needs to make the *end-to-end* part visible immediately (e.g., a one-line pipeline: "model → retrieval → API → e ↓ > deployed"), not just assert competence.

Otherwise it's a headline, not proof.

Pressure-test output

AI Engineering Portfolio Build / Portfolio sitemap pressure-test against proof statement ▾

About — Earns its place, but keep it short and don't let it re-argue the claim in prose. Bio + focus areas, not a second pitch.

Contact — Earns it by definition; it's your one action.

2. What's missing or unnecessary

Missing: explicit *evidence per component*. Your proof statement makes five specific claims. Right now nothing in the sitemap guarantees a visitor sees evidence of "evaluation" or "production deployment" — those are the two easiest parts of the claim to fake and the two recruiters will be most skeptical of. Each Selected Work entry needs a visible, specific artifact for these: eval methodology/metrics, and deployment reality (live endpoint, users, uptime, latency — something concrete, not "deployed to production" as a phrase).

Unnecessary (as currently framed): Webometrics research, unless you can reframe it to show a shipped system with backend/eval/deployment components. If it's just research output, it belongs in About as a credential line, not as a Selected Work case study competing for attention against your strongest proof point.

3. Reordering

Hero → Selected Work → About → Contact is the right order for this audience — recruiters want proof before biography. Don't change the section order. Do put your strongest end-to-end proof (RAG/multi-agent) *first* within Selected Work, not last — right now it's listed third, which buries your best evidence under two weaker pieces.

4. The single most important change

Your proof statement makes five claims, but only one of your three projects can defend all five. Right now you're relying on the RAG/multi-agent project to single-handedly prove your entire pitch, while the other two projects mostly support "I can do ML," which is a narrower and less differentiated claim.

Fix: either (a) find or build a second project that visibly demonstrates the backend API + evaluation + deployment chain, so the claim isn't resting on one example, or (b) restructure each existing project's write-up to explicitly show which parts of the pipeline it covers — and cut or demote whichever project can't show at least backend/eval/deployment evidence, even if it's technically impressive. A recruiter skimming for 30 seconds should be able to see the full pipeline proven at least twice, not asserted once and implied by association twice more.



Pressure-test output - continued

AI Engineering Portfolio Build / Portfolio sitemap pressure-test against proof statement ▾

Here's the pressure test — not a rubber stamp.

1. Does every section earn its place?

Hero — Earns it, conditionally. "Clearly state my AI engineering claim" is right, but a claim alone doesn't do work. The hero needs to make the *end-to-end* part visible immediately (e.g., a one-line pipeline: "model → retrieval → API → eval → deployed"), not just assert competence. Otherwise it's a headline, not proof.

Selected Work — This is where your claim lives or dies, and right now it's uneven:

- **Multi-Agent/RAG systems:** strong match. This is likely the only project that can touch all five components in your proof statement (model/retrieval design, backend APIs, evaluation, deployment).
- **TensorFlow open-source contribution:** proves ML depth and code quality, but does it show *you* shipping a system end-to-end, or contributing to someone else's? If it's a PR/patch, it supports "model design" credibility but says nothing about retrieval, APIs, evaluation, or deployment. Right now it's doing decoration work, not proof work.
- **Webometrics ranking research:** this is the weakest fit. "Research" and "ranking" signal analysis, not a *shipped AI system*. Unless this project involved you building and deploying something (a pipeline, a service, a model in production), it actively dilutes the proof statement rather than supporting it.

About — Earns its place, but keep it short and don't let it re-argue the claim in prose. Bio + focus areas, not a second pitch.

Contact — Earns it by definition; it's your one action.

2. What's missing or unnecessary

Missing: explicit *evidence per component*. You ▾ of statement makes five specific claims. Right now nothing in the sitemap guarantees a visitor sees evidence of "evaluation" or